

Workshop: High-performance computing for economists

Lars Vilhuber¹ John M. Abowd¹ Richard Mansfield¹
Hautahi Kingi¹ Flavio Stanchi¹ Sida Peng¹
Kevin L. McKinney


¹Cornell University, Economics Department,

August 17-19, 2015

Workshop: High-performance computing for economists


HPC

Back in the days...



HPC


Back in the days...



RAM: 2,000 words (2kB); Speed: 2 MHz

Source: Wikipedia

They went to the moon



Source: Flickr

Big progress



RAM: 2×32 kB; Speed: 1 MHz, \$1,500 (today's USD)

Wikipedia

Today



RAM: 2×1024^2 kB; Speed: $1.700 \text{ MHz} \times 4$
\$700 (today's USD) Source: Wikipedia

We still fly to the moon



Source CNET

This is where you can go

Stampede (no. 6 on Top500 as of June 2013)



This is where you can go

Stampede (no. 6 on Top500 as of June 2013)



RAM: 192×1024^3 kB, Speed: 2,700 Mhz \times 462,462

Source: TACC

But first...

But first...



<http://viewfromwitsend.wordpress.com/>

What do you learn in a Ph.D. program?

What do you learn in a Ph.D. program?

How to learn...

Goal of this class

Goal of this class

To open new doors, to be able to conceive of problems that you didn't think had a feasible solution.

Goal of this class

To open new doors, to be able to conceive of problems that you didn't think had a feasible solution.

To broaden your knowledge about what you do NOT know

So in order to do that...

So in order to do that...



So in order to do that...



Overview

Day 1

- ▶ Programming basics (Lars)

Overview

Day 1

- ▶ Programming basics (Lars)
 - ▶ Choosing an editor

Overview

Day 1

- ▶ Programming basics (Lars)
 - ▶ Choosing an editor
 - ▶ How to structure programs, texts, etc.

Overview

Day 1

- ▶ Programming basics (Lars)
 - ▶ Choosing an editor
 - ▶ How to structure programs, texts, etc.
 - ▶ A clean sequence of programs

Overview

Day 1

- ▶ Programming basics (Lars)
 - ▶ Choosing an editor
 - ▶ How to structure programs, texts, etc.
 - ▶ A clean sequence of programs
 - ▶ NX, SSH, Linux, request an account on cluster

Overview

Day 1

- ▶ Programming basics (Lars)
 - ▶ Choosing an editor
 - ▶ How to structure programs, texts, etc.
 - ▶ A clean sequence of programs
 - ▶ NX, SSH, Linux, request an account on cluster
 - ▶ Basic scripting

Overview

Day 1

- ▶ Programming basics (Lars)
 - ▶ Choosing an editor
 - ▶ How to structure programs, texts, etc.
 - ▶ A clean sequence of programs
 - ▶ NX, SSH, Linux, request an account on cluster
 - ▶ Basic scripting
- ▶ Basics of version control (Lars)

Overview

Day 1

- ▶ Programming basics (Lars)
 - ▶ Choosing an editor
 - ▶ How to structure programs, texts, etc.
 - ▶ A clean sequence of programs
 - ▶ NX, SSH, Linux, request an account on cluster
 - ▶ Basic scripting
- ▶ Basics of version control (Lars)
 - ▶ File-system based version control

Overview

Day 1

- ▶ Programming basics (Lars)
 - ▶ Choosing an editor
 - ▶ How to structure programs, texts, etc.
 - ▶ A clean sequence of programs
 - ▶ NX, SSH, Linux, request an account on cluster
 - ▶ Basic scripting
- ▶ Basics of version control (Lars)
 - ▶ File-system based version control
 - ▶ More formal version control (Subversion, Git)

Overview

Day 1

- ▶ Programming basics (Lars)
 - ▶ Choosing an editor
 - ▶ How to structure programs, texts, etc.
 - ▶ A clean sequence of programs
 - ▶ NX, SSH, Linux, request an account on cluster
 - ▶ Basic scripting
- ▶ Basics of version control (Lars)
 - ▶ File-system based version control
 - ▶ More formal version control (Subversion, Git)
 - ▶ Working with servers

Overview

Day 1

- ▶ Programming basics (Lars)
 - ▶ Choosing an editor
 - ▶ How to structure programs, texts, etc.
 - ▶ A clean sequence of programs
 - ▶ NX, SSH, Linux, request an account on cluster
 - ▶ Basic scripting
- ▶ Basics of version control (Lars)
 - ▶ File-system based version control
 - ▶ More formal version control (Subversion, Git)
 - ▶ Working with servers
 - ▶ Setting up infrastructure at Cornell

Overview

Day 1

- ▶ Programming basics (Lars)
 - ▶ Choosing an editor
 - ▶ How to structure programs, texts, etc.
 - ▶ A clean sequence of programs
 - ▶ NX, SSH, Linux, request an account on cluster
 - ▶ Basic scripting
- ▶ Basics of version control (Lars)
 - ▶ File-system based version control
 - ▶ More formal version control (Subversion, Git)
 - ▶ Working with servers
 - ▶ Setting up infrastructure at Cornell
- ▶ HP resources at Cornell, elsewhere

Overview

Overview

Day 3

Structure of the class

Teaching...

We'll take you on a 4,000 m flight through topics...

Structure of the class

Teaching...

We'll take you on a 4,000 m flight through topics...

... and practice

... and then swoop in on some examples, leaving ample time to practice it.

Choosing editors

Why does choosing editors matter?

The (applied) research process iterates through writing papers and doing estimation. You want to use the appropriate tools for each task.

Integrated or separate

- ▶ You can use native tools that come with each word processing facility/programming language/etc.
- ▶ Not all of them will have one.
- ▶ Not all of them will work on all platforms.
- ▶ You will likely use multiple tools

Choosing an editor

... or system

Separate editors and systems

- ▶ MS Word and math editor (Windows/OSX but compatibility issues)
- ▶ LibreOffice (Windows/OSX/Linux but not as good)
- ▶ Notepad++ (Windows)
- ▶ Gedit, (X)Emacs, Kate (Linux)
- ▶ Sublime Text (OSX)
- ▶ Atom (all, see also MS Visual Studio Code)

L^AT_EX: all platforms, but some GUIs are not cross-platform, ease of use varies:

- ▶ TeXstudio (all platforms)
- ▶ TeXMaker (all platforms)
- ▶ Scientific Workplace (Windows, mythical Linux)
- ▶ TeXWorks+Miktex
- ▶ TEXnicCenter
- ▶ and (many more)

Choosing an editor

... or system

Integrating programming and running

- ▶ IDE (Eclipse, ActiveState Komodo, etc.)
- ▶ Native programming GUIs (SAS, Matlab, Stata)
- ▶ Gedit, (X)Emacs (with add-on functionality)

Integrating programs and text/results

- ▶ SWeave/knitr (integrates \LaTeX and R)
- ▶ RStudio (GUI to R and SWeave/knitr)
- ▶ Shiny (web interface to R with dynamic results)
- ▶ StatRep (Integrated SAS and \LaTeX , Source 1, Source 2)

Structuring programs

Structuring programs

Easy...

Listing 1: mystuff.sas

```
1 data "C:\Users\Me\CensusChina.sas7bdat";  
2 set "C:\Users\Me\CensusChina.sas7bdat";  
3 earn=log(earn);  
4 run;  
5 proc reg data="C:\Users\Me\CensusChina.sas7bdat";  
6 model earn = sex education experience;  
7 run;
```

What can possibly be wrong about that?

Structuring programs 2

Easier...

Listing 2: mystuff.do

```
1 use "C:\Users\Me\CensusChina.dta"  
2 replace earn=log(earn)  
3 regress earn sex education experience  
4 save, replace
```

What can possibly be wrong about that?

Structuring programs 3

Actually...

Everything!

- ▶ Name of program: uninformative
- ▶ Destruction of original data: program cannot be re-run for same results
- ▶ No portability: cannot be run anywhere else
- ▶ No explanation: why are we doing this?

But of course, nobody does that, right?

Structuring programs 4

Better...?

Listing 3: china-regression.sas

```
1 data logCensusChina;  
2     set "C:\Users\Me\CensusChina.sas7bdat";  
3     earn=log(earn);  
4 run;  
5 proc reg data=logCensusChina;  
6 model earn = sex education experience;  
7 run;
```

Structuring programs 4

Better...?

Listing 4: china-regression.sas

```
1 data logCensusChina;  
2     set "C:\Users\Me\CensusChina.sas7bdat";  
3     earn=log(earn);  
4 run;  
5 proc reg data=logCensusChina;  
6 model earn = sex education experience;  
7 run;
```

Somewhat...

Structuring programs 5

Addressing these issues

- ▶ Naming of programs: [here](#)
- ▶ Commenting: [here](#)
- ▶ Versioning: [up next](#)
- ▶ Portability and Data management: [tomorrow](#)

Key notions about naming

Think of yourself as highly amnesiac...

- ▶ The research paper you are writing now will be submitted, rejected, worked on, questioned...

Key notions about naming

Think of yourself as highly amnesiac...

- ▶ The research paper you are writing now will be submitted, rejected, worked on, questioned...
- ▶ ... by others and yourself

Key notions about naming

Think of yourself as highly amnesiac...

- ▶ The research paper you are writing now will be submitted, rejected, worked on, questioned...
- ▶ ... by others and yourself
- ▶ ... in intervals of weeks, months, years...

Key notions about naming

Think of yourself as highly amnesiac...

- ▶ The research paper you are writing now will be submitted, rejected, worked on, questioned...
- ▶ ... by others and yourself
- ▶ ... in intervals of weeks, months, years...
- ▶ Your future research assistant and the future YOU will need to understand how to go through it.

Naming

The really bad

mystuff.R

read.R

version2.R

ols.sas

Naming

The really bad

mystuff.R

read.R

version2.R

ols.sas

The bad

readCensus.R

readBLS.R

prepareCensus.R

runOLS.sas

Naming

Better

01_readBLS.R

02_readCensus.R

03_prepareCensus.R

04_create_analysis_data.R

05_runOLS.sas

Naming

Better

```
01_readBLS.R  
02_readCensus.R  
03_prepareCensus.R  
04_create_analysis_data.R  
05_runOLS.sas
```

Even better

```
01_01_readBLS.R  
02_01_readCensus.R  
02_02_prepareCensus.R  
03_01_create_analysis_data.R  
04_01_runOLS.sas  
README.txt
```

Naming

Going overboard?

```
icf/ctrlprogs/control_icf.sas
icf/ctrlprogs/parameters_icf.sas
icf/library/macros/icf_cleanup.sas
icf/library/macros/icf_impute_county_res.sas
icf/library/macros/licf_findnum.sas
icf/library/macros/licf_proxy.sas
icf/library/macros/licf_stars1.sas
icf/library/macros/licf_tgrlatlongs.sas
icf/library/sasprogs/01_icfqa.sas
icf/library/sasprogs/01_icf.sas
icf/library/sasprogs/02_icfqa.sas
icf/library/sasprogs/02_icf.sas
icf/library/sasprogs/03_icfqa.sas
icf/library/sasprogs/03_icf.sas
[snip]
icf/library/sasprogs/19_icf.sas
```

Naming

Going overboard?

```
icf / ctrlprogs / control_icf .sas  
icf / ctrlprogs / parameters_icf .sas  
icf / library / macros / icf_cleanup .sas  
icf / library / macros / icf_impute_county_res .sas  
icf / library / macros / licf_findnum .sas  
icf / library / macros / licf_proxy .sas  
icf / library / macros / licf_stars1 .sas  
icf / library / macros / licf_tgrlatlongs .sas  
icf / library / sasprogs / 01_icfqa .sas  
icf / library / sasprogs / 01_icf .sas  
icf / library / sasprogs / 02_icfqa .sas  
icf / library / sasprogs / 02_icf .sas  
icf / library / sasprogs / 03_icfqa .sas  
icf / library / sasprogs / 03_icf .sas  
[snip]  
icf / library / sasprogs / 19_icf .sas
```

```
ehf / ctrlprogs / control_ehf .sas  
ehf / library / macros / read_bls .sas  
ehf / library / sasprogs / 01_ehf .sas  
[snip]
```


Naming

With minor modification

```
icf /ctrlprogs/control_icf.sas
icf /ctrlprogs/parameters_icf.sas
icf /library/macros/icf_cleanup.sas
icf /library/macros/icf_impute_county_res.sas
icf /library/macros/licf_findnum.sas
icf /library/macros/licf_proxy.sas
icf /library/macros/licf_stars1.sas
icf /library/macros/licf_tgrlatlongs.sas
icf /library/sasprogs/01_icf.sas
icf /library/sasprogs/02_icf.sas
icf /library/sasprogs/03_icf.sas
[snip]
icf /library/sasprogs/19_icf.sas
icf /library/sasprogs/01_icfqa.sas
icf /library/sasprogs/02_icfqa.sas
icf /library/sasprogs/03_icfqa.sas
```

Can you figure out in what sequence to run them?

Why SSH?

Most compute clusters have ONLY SSH access

It is thus worthwhile to learn enough about it here, in order to be functional there: CAC “Red Cloud”, Amazon Cloud, XSEDE.

Linux rules... the HPC world

All 10 of the top 10 TOP500 computers run Linux (as the compiler front-end, if not compute OS)

Graphical access

Two types of graphical access

- ▶ with an “X server” (native in Linux, optional in Windows and OSX)

Graphical access

Two types of graphical access

- ▶ with an “X server” (native in Linux, optional in Windows and OSX) → standard way on most clusters

Graphical access

Two types of graphical access

- ▶ with an “X server” (native in Linux, optional in Windows and OSX) → **standard way on most clusters**
- ▶ using NX client software for improved experience

Basic Linux, basic scripting

Why worry?

You will end up doing something on the command line

- ▶ Launch a program from a compute-cluster job

Why worry?

You will end up doing something on the command line

- ▶ Launch a program from a compute-cluster job
- ▶ Launch a job submission

Why worry?

You will end up doing something on the command line

- ▶ Launch a program from a compute-cluster job
- ▶ Launch a job submission
- ▶ Basic scripting

Linux in 2 minutes

- ▶ ls - will list the contents of a directory
- ▶ cd - will “change directory”
- ▶ cd .. (note the spaces) will go up a directory
- ▶ cd (name) will go into the directory (name)
- ▶ rm (name) will delete
- ▶ mkdir (name) will create a directory called (name)
- ▶ vi (name) will open a venerable command line editor for file (name)

Linux in 2 minutes

- ▶ ls - will list the contents of a directory
- ▶ cd - will “change directory”
- ▶ cd .. (note the spaces) will go up a directory
- ▶ cd (name) will go into the directory (name)
- ▶ rm (name) will delete
- ▶ mkdir (name) will create a directory called (name)
- ▶ vi (name) will open a venerable command line editor for file (name) (CAUTION: to exit, hit ESC, then :q!)

Basic scripting in Linux

A basic loop on the command line

```
1  for (( i; i<10; i++ ))  
2  do  
3    echo $i  
4  done  
5  for i in 1 3 7 99  
6  do  
7    echo $i  
8  done
```

Source: [1]

Capturing output

You can capture the output from a command

```
> seq 1 3  
1  
2  
3
```

Now let's use that:

```
for i in $(seq 1 3)  
do  
    echo $i  
done
```

Basic scripting in Linux

Use for practical things

Remember that ICF program sequence? How would we go about starting 19 programs in sequence?

```
for program in $(ls *_icf.sas)
do
    sas $program
done
```

Advanced linux in 2 minutes

The gateway to everything

man

or try <http://www.linuxmanpages.com> or <http://linux.die.net/man/>

The toolkit

- ▶ sed
- ▶ grep
- ▶ awk
- ▶ regex (regular expressions)

Advanced scripting in Linux

Use for practical things

What if I'm running 100s of programs, and trying to figure out if any of them have errors?

```
for logfiles in $(ls *_icf.log)
do
    grep ERROR $logfiles
done
```


Now let's try it out

Next section

Next section