

# Workshop: High-performance computing for economists

Lars Vilhuber<sup>1</sup> John M. Abowd<sup>1</sup> Richard Mansfield<sup>1</sup>  
Hautahi Kingi<sup>1</sup> Flavio Stanchi<sup>1</sup> Sida Peng<sup>1</sup>  
Kevin L. McKinney

<sup>1</sup>Cornell University, Economics Department,

August 17-19, 2015

# Basic subroutine programming

## Goal

- ▶ Show the basics of proper subroutine programming
- ▶ Advantages, pitfalls
- ▶ Examples in R
- ▶ Later: generalization and differences in other programming languages

# Control structures in programming languages

## Mostly generic

- ▶ `if, else`: testing a condition [R, SAS]
- ▶ `for`: execute a loop a fixed number of times [R, in SAS: `do`]
- ▶ `while`: execute a loop while a condition is true [R,SAS]
- ▶ `until`: execute a loop until a condition is true [SAS]
- ▶ `repeat`: execute an infinite loop [R]
- ▶ `break`: break the execution of a loop [R, SAS]
- ▶ `next`: skip an iteration of a loop [R]
- ▶ `return`: exit a function [R]

# Control structures: if

# Control structures: if

## ... in R

```
1  if(<condition>) {  
2  ## do something  
3  } else {  
4  ## do something else  
5  }  
6  if(<condition1>) {  
7  ## do something  
8  } else if(<condition2>) {  
9  ## do something different  
10 } else {  
11 ## do something different  
12 }
```

# Control structures: if

## ... in R

```

1  if(<condition>) {
2  ## do something
3  } else {
4  ## do something else
5  }
6  if(<condition1>) {
7  ## do something
8  } else if(<condition2>) {
9  ## do something different
10 } else {
11 ## do something different
12 }

```

## ... in SAS

```

1  if (<condition>) then do;
2  ## do something
3  end; else do;
4  ## do something else
5  end;
6  if (<condition1>) then do;
7  ## do something
8  else if (<condition2>) then do;
9  ## do something different
10 end; else do;
11 ## do something different
12 end;

```

# Control structures: for

Run through a fixed sequence of numbers (or in R, a sequence of vectors)

# Control structures: for

Run through a fixed sequence of numbers (or in R, a sequence of vectors)

## simple loop in R

```
1  for(i in 1:10) {  
2    print(i)  
3  }
```



# Control structures: for

Run through a fixed sequence of numbers (or in R, a sequence of vectors)

## simple loop in R

```
1  for(i in 1:10) {  
2  print(i)  
3  }
```

## ... in SAS

```
1  do i = 1 to 10;  
2  put i ;  
3  end;
```

# Control structures: for

Across programming languages, some flexibility:

# Control structures: for

Across programming languages, some flexibility:

## Equivalent loops in R

```
1 x <- c("a", "b", "c", "d")
2 for(i in 1:4) {
3   print(x[i])
4 }
5 for(i in x) {
6   print(i)
7 }
8 for(i in 1:4) print(x[i])
```

## ... in SAS

```
1 do i = 1 to 10;
2   put i;
3 end;
```

# Interrupting loops

## How a loop ends

- ▶ at the end
- ▶ by resetting the counter to the end value, or by setting the looping condition to its exit value explicitly
- ▶ when it encounters a `break` (for `repeat` loops in R)

# What do we loop over?

Within the loop, something is done

```
1 for (i in 1:100000) {  
2  ## do something here  
3  # stuff (1 line)  
4  # stuff (2nd line)  
5  # done  
6 }
```

# What do we loop over?

Within the loop, something is done

```
1 for (i in 1:100000) {  
2  ## do something here  
3  # stuff (1 line)  
4  # stuff (2nd line)  
5  # done  
6 }
```

What if that is really complicated?

```
1 for (i in 1:100000) {  
2  ## do something really complicated here (398 lines)  
3 }
```

# Subroutines

## Breaking into discrete chunks

```
1  for (i in 1:100000) {  
2    ## do something really complicated here  
3    ## (398 lines)  
4  }
```

# Subroutines

## Breaking into discrete chunks

```
1  for (i in 1:100000) {  
2  ## do something really complicated here  
3  ## (398 lines)  
4  }
```

In the previous example, this was unwieldy (which loop are you closing on line 400?)



# Subroutines

## Breaking into discrete chunks

```
1  for (i in 1:100000) {  
2    ## do something really complicated here  
3    ## (398 lines)  
4  }
```

In the previous example, this was unwieldy (which loop are you closing on line 400?)

We can start by breaking out the complicated stuff into a well-defined subroutine.

# Subroutines

## Breaking into discrete chunks

```

1  for (i in 1:100000) {
2    ## do something really complicated here
3    ## (398 lines)
4  }

```

We can start by breaking out the complicated stuff into a well-defined subroutine.

In the previous example, this was unwieldy (which loop are you closing on line 400?)

```

1  for (i in 1:100000) {
2    ## new subroutine
3    do_something_complicated(args=something)
4    ##
5  }

```

# Subroutines, procedures, etc.

## What do you call a subroutine?

“The same things in different [programming] languages can have different names. Programs, Procedures, Functions, **Subroutines**, Subprograms, Subqueries ... these words all have very similar meanings. [...] We can call an object, it executes and performs a complex process. Whether that object is called any one of the above list depends on what programming language it is written in, whether a human being can call it, or whether it has to be called by another program or one of the other names on the list.”

Source

# An practical overview of subroutine programming

## We will use R...

... but expand to cover what this might look like in SAS (macros), Stata (programs, ado files), Matlab (functions) tomorrow.

... use this as a building block to scale to HPC.

## You already use them!

Much functionality in SAS, Stata, Matlab, R is implemented as a subroutine:

- ▶ `proc import` in SAS (compare call to log file)
- ▶ Most statistical functions in Stata (`regress` used earlier is actually in `../ado/base/r/regress.ado`)

You can create them too!

# Functions in R

## R objects of class "function"

```
1 f <- function(<arguments>) {  
2   ## do something here  
3 }
```

- ▶ Functions can be arguments to other functions
- ▶ Functions can be nested (even recursive)
- ▶ Functions can have named arguments with default values, or missing values

We draw on Peng's "Computing for Data Analysis, Week 2" for this section.

# Basic examples

## Let's start with a simple example

```
1  # This is a very simple function
2  # It simply draws a draws from a normal with mean=b
3  f <- function(a=1,b=0) {
4      rnorm(a,mean = b)
5  }
```

# Basic examples

## Let's start with a simple example

```
1  # This is a very simple function
2  # It simply draws a draws from a normal with mean=b
3  f <- function(a=1,b=0) {
4      rnorm(a,mean = b)
5  }
```

The function is defined, has two arguments with default values, and a really short name.



# Basic examples

## Let's start with a simple example

```

1  # This is a very simple function
2  # It simply draws a draws from a normal with mean=b
3  f <- function(a=1,b=0) {
4      rnorm(a,mean = b)
5  }

```

The function is defined, has two arguments with default values, and a really short name.

## Results

```

1  > source("3-1-simple-function.R")
2  > f
3  function(a=1,b=0) {
4      rnorm(a,mean = b)
5  }
6  > f()
7  [1] 0.1709822
8  > set.seed(10)
9  > f(a=10,b=5)
10 [1] 5.018746 4.815747 3.628669 4.400832 5.294545 5.389794 3.791924 4.636324
11 [9] 3.373327 4.743522

```

# Defining function arguments

Functions have *named* arguments which potentially have *default* values.

- ▶ The formal arguments are the arguments included in the function definition
- ▶ Not every function call in R makes use of all the formal arguments
- ▶ Function arguments can be missing, or can have default values

# Function arguments at invocation

Function arguments can be passed by position (*positional*) or by explicit name (in which case, position doesn't matter).

## Listing 1: Equivalent calls

```
1 > set.seed(10)
2 > mean( f (b=5,a=10))
3 [1] 4.509343
4 > set.seed(10)
5 > mean( f (a=10,b=5))
6 [1] 4.509343
7 > set.seed(10)
8 > mean( f (10,5))
9 [1] 4.509343
```

However, it is **good practice** to use *named* arguments, as they make use of the function (especially when more complex) easier and more transparent.

# Lazy evaluation

## Arguments can remain unused:

```
f <- function(a, b) {  
  a^2  
}  
> f(2)  
[1] 4
```

b was never used.

# Lazy evaluation

## Arguments can remain unused:

```
f <- function(a, b) {
  a^2
}
> f(2)
[1] 4
```

`b` was never used.

## Arguments can be faulty, but don't lead to problems until used

```
f <- function(a, b) {
  print(a)
  print(b)
}
> f(45)  # note no value provided for b
[1] 45
Error in print(b) : argument "b" is missing, with no default
>
```

The absence of a value for `b` only led to an error at the time it was called.

# Scoping

## Scope of functions

```
f <- function(a,b) {  
  print(a)  
  print(c)  
  c <- paste(a,b)  
  print (c)  
}
```

Try it out:

# Scoping

## Scope of functions

```
f <- function(a,b) {
  print(a)
  print(c)
  c <- paste(a,b)
  print (c)
}
```

Try it out:

```
1 > f(c('a'),c('b'))
2 [1] "a"
3 function (... , recursive = FALSE) .Primitive("c") <== ERROR!
4 [1] "a_b"
```

# Scoping

## What happened?

`c` was not defined, leading to the error on line 4. Line 5 reports what happens to `c` once it is defined.

## Understanding the scope

```
5 > c <- c( 'Nothing ' )
6 > f( c( 'a' ), c( 'b' ) )
7 [1] "a"
8 [1] "Nothing"
9 [1] "a_b"
```

So what value does `c` now have?



# Scoping

## What happened?

`c` was not defined, leading to the error on line 4. Line 5 reports what happens to `c` once it is defined.

## Understanding the scope

```
12 > c <- c( 'Nothing' )
13 > f(c( 'a' ), c( 'b' ))
14 [1] "a"
15 [1] "Nothing"
16 [1] "a_b"
```

So what value does `c` now have?

```
17 > print(c)
18 [1] "Nothing"
```

# Scoping

## Understanding the scope of variables

- ▶ In the example, `c` was defined both inside the function and outside.

# Scoping

## Understanding the scope of variables

- ▶ In the example, `c` was defined both inside the function and outside.
- ▶ R took the “global” value of `c` ...

# Scoping

## Understanding the scope of variables

- ▶ In the example, `c` was defined both inside the function and outside.
- ▶ R took the “global” value of `c` ...
- ▶ until the function defined it, at which point it took a new value internal to the function

# Scoping

## Understanding the scope of variables

- ▶ In the example, `c` was defined both inside the function and outside.
- ▶ R took the “global” value of `c` ...
- ▶ until the function defined it, at which point it took a new value internal to the function
- ▶ ... which was only used until the function ended. The “global” value had not changed.

# Scoping

## Understanding the scope of variables

- ▶ In the example, `c` was defined both inside the function and outside.
- ▶ R took the “global” value of `c` ...
- ▶ until the function defined it, at which point it took a new value internal to the function
- ▶ ... which was only used until the function ended. The “global” value had not changed.

# Scoping

## Understanding the scope of variables

- ▶ In the example, `c` was defined both inside the function and outside.
- ▶ R took the “global” value of `c` ...
- ▶ until the function defined it, at which point it took a new value internal to the function
- ▶ ... which was only used until the function ended. The “global” value had not changed.

## Other programming languages

Each programming language has a different way of handling these. Take care to read up on it.

# Searching for variables (and functions)

- ▶ The global environment or the user's workspace is always the first element of the search list and the **base** package is always the last.
- ▶ The order of the packages on the search list matters!
- ▶ User's can configure the order of packages as they get loaded on startup ...
- ▶ When a user loads a package with library the namespace of that package gets put in position 2 of the search list (by default) and everything else gets shifted down the list.



# Scoping (name resolution) rules

## Lexical scoping

**Lexical** (or **static**) resolution can be determined at compile time, and is also known as **early binding**

## Dynamic scoping

**Dynamic** resolution can in general only be determined at run time, and thus is known as **late binding**.

Most modern languages use lexical scoping for variables and functions, though de facto dynamic scoping is common in macro languages, which do not directly do name resolution. [1]

# Scoping rules

## Back to R

```
f <- function(a,b) {  
  print(paste(a,b))  
  print(c)  
}
```

- ▶ Named (formal) arguments ( $a, b$ ) are always local
- ▶ Free variables (not defined in the call) ( $c$ ) are subject to scoping rules

# Lexical scoping in R

Lexical scoping in R means that

*the values of free variables are searched for in the environment in which the function was defined.*

What is an environment?

- ▶ *An environment is a collection of (symbol, value) pairs, i.e.  $x$  is a symbol and 3.14 might be its value.*
- ▶ *Every environment has a parent environment; it is possible for an environment to have multiple children?*
- ▶ *A function + an environment = a closure or function closure.*

# Safe use of scoping

## Define all variables as arguments

To avoid confusion about where variables are defined, unless you have a good reason to deviate from this

- ▶ Define all variables as arguments to the function
- ▶ In other languages, explicitly define the scope of a variable

# Scoping in SAS

## Scope of macro variables in SAS

### Define a similar program in SAS

```
%macro myprogram(a=,b=);  
%let c=&a.&b.;  
%put &c.;  
%mend;
```

# Scoping in SAS

## Scope of macro variables in SAS

Define a similar program in SAS

```
%macro myprogram(a=,b=);  
%let c=&a.&b.;  
%put &c.;  
%mend;
```

Call it and assess where  $c$  comes from:

```
%myprogram(a=a,b=b);  
%put &c.;  
%let c=Nothing;%put &c.;  
%myprogram(a=a,b=b);  
%put &c.;
```

# Scoping in SAS

## with results

```

1  %macro myprogram(a=,b=);
2  %let c=&a.&b.;
3  %put &c.;
4  %mend;
6  %myprogram(a=a,b=b);
ab
WARNING: Apparent symbolic reference C not resolved.
7  %put &c.;
&c.
8  %let c=Nothing;%put &c.;
Nothing
9  %myprogram(a=a,b=b);
ab
10 %put &c.;
ab

```

# Scoping in SAS

Better:

```
1  %macro myprogram(a=,b=);  
2  %local c;  
3  %let c=&a.&b.;  
4  %put &c.;  
5  %mend;  
6  %let c=Nothing;  
7  %myprogram(a=a,b=b);  
ab  
8  %put &c.;  
Nothing
```



# Efficient use of scoping

## Differences that need to be taken into account

- ▶ Scoping rules can be leveraged to improve optimization (see Peng's Coursera course and others)
- ▶ The use of scoping differs across languages (what is feasible in R cannot be simply translated into Java or SAS or Stata)

# Naming rules

## Naming is important

- Naming is important - both of functions and of arguments:  
Compare the following two functions:

```

1  f <- function(a,b,c) {
2    x <- sample(y,c)
3    lm(a ~ b , data=x )
4  }
```

and

```

1  sample_reg <- function(lhs,rhs,samplesize=10,data=) {
2    subset <- sample(data,samplesize)
3    lm(lhs ~ rhs , data=subset )
4  }
```

- Give functions and variables meaningful names

# Documentation

## Provide some source-code level documentation

```

1  # Author: Lars Vilhuber
2  # This program defines a regression on a sample.
3  #
4  # Usage: sample_reg(lhs, (list of rhs), samplesize=10, data=data frame)
5  #
6  # Samplesize defaults to 10 observations, and will sample without replacement
7  # Regression is simple linear regression with no options
8  # Function returns an object of type lm
9
10 sample_reg <- function(lhs, rhs, samplesize=10, data=) {
11   subset <- sample(data, samplesize)
12   lm(lhs ~ rhs, data=subset)
13 }

```

- More sophisticated documentation: see package documentation

# Naming rules

## Robustness is important

- Think about backward compatibility (and your program library). Say you first used this:

```

1  sample_reg <- function(lhs,rhs,samplesize=10,data=) {
2      subset <- sample(data,samplesize)
3      lm(lhs ~ rhs , data=subset )
4  }
5  sample_reg(earnings , education ,5 ,data=cps)

```

# Naming rules

## Robustness is important

- Think about backward compatibility (and your program library). Say you first used this:

```

1  sample_reg <- function(lhs,rhs,samplesize=10,data=) {
2      subset <- sample(data,samplesize)
3      lm(lhs ~ rhs , data=subset )
4  }
5  sample_reg(earnings,education,5,data=cps)

```

- Now you extend your model

```

1  sample_reg <- function(lhs,rhs,samplesize=10,data=) {
2      subset <- sample(data,samplesize)
3      lm(log(lhs) ~ rhs , data=subset )
4  }
5  sample_reg(earnings,education,5,data=cps)

```

# Naming rules

## Robustness is important

- ▶ You should either give your function a different name, or make the call robust to both the “old” way and the “new” way:

```

1  sample_reg <- function(lhs,rhs,samplesize=10,data=,logs = false ) {
2    subset <- sample(data,samplesize)
3    if ( logs ) {
4      _lhs <- log(lhs)
5    }
6    else {
7      _lhs <- lhs
8    }
9    lm(_lhs ~ rhs , data=subset )
10 }
```

which will work for both calls...

## The power of functions

# The power of functions

## Why bother with functions?

- ▶ Initial example: putting 398 command lines into separate file (ease of use)
- ▶ Expansion on that: re-using the function across multiple projects (function library)
- ▶ Your function is a complete specification. You know want to vary or perturb all 25 parameters slightly, for robustness checks.

```
1  for (i in 1:1000) {  
2    for (j in 1:1000) {  
3      my_regression(model=base, xi=i, xj=j)  
4    }  
5  }
```



# The power of functions

## Why bother with functions?

- ▶ You want to database the results:

```
1  for (i in 1:1000) {  
2    for (j in 1:1000) {  
3      results_db[i,j]=my_regression(model=base, xi=i, xj=j)  
4    }  
5  }
```

# The power of functions

## The ultimate power of functions: scaling

- You want to speed the whole thing up by parallelizing:

```

1  library(doMC)
2  registerDoMC()
3  foreach (i = 1:1000, .combine=cbind) %dopar% {
4    for (j in 1:1000) {
5      results_db[i,j]=my_regression(model=base, xi=i, xj=j)
6    }
7  }

```

which will run 1000 parallel threads, on as many cores as you can. (see doMC vignette or the help in your Rstudio installation)

# Take-away

## Major points

- ▶ Subroutines are a powerful tool to write clean, understandable code
- ▶ Subroutines (functions, macros, programs) are present in some form in all statistical programming languages
- ▶ Use consistent, clear naming
- ▶ Use robust subroutines, expanding them into a library that you can use and share across projects
- ▶ Clean subroutines are a critical component to scaling your analysis (parallelization)

# Take-away

## Additional items

- ▶ learn how to debug (different in each language, also critical to scaling)
- ▶ subroutines don't magically make your code efficient - they allow you to figure out which portions are not

# Take-away

## Additional items

- ▶ consider publishing your routines
  - ▶ R packages `http://r-pkgs.had.co.nz/` with publication in R Journal
  - ▶ Stata packages (user-contributed commands, ado files) from the Boston College Statistical Software Components (SSC) archive with publication in Stata Journal
  - ▶ Journal of Statistical Software
  - ▶ Github... (Stata, R devtools )